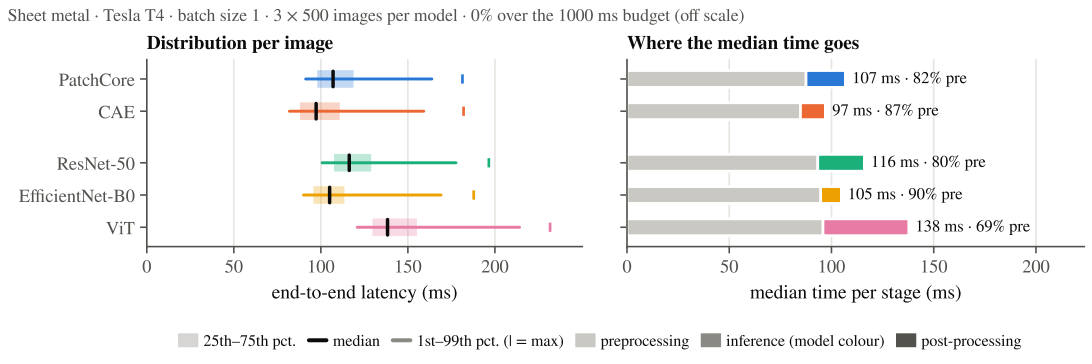
Sheet metal · Tesla T4 · batch size 1 · 3 × 500 images per model · 0% over the 1000 ms budget (off scale)
Distribution per image
PatchCore
CAE
ResNet-50
EfficientNet-B0
ViT
0
50
100
150
200
end-to-end latency (ms)
Where the median time goes
107 ms · 82% pre
97 ms · 87% pre
116 ms · 80% pre
105 ms · 90% pre
138 ms · 69% pre
0
50
100
150
200
median time per stage (ms)
25th–75th pct.
median
1st–99th pct. (| = max)
preprocessing
inference (model colour)
post-processing