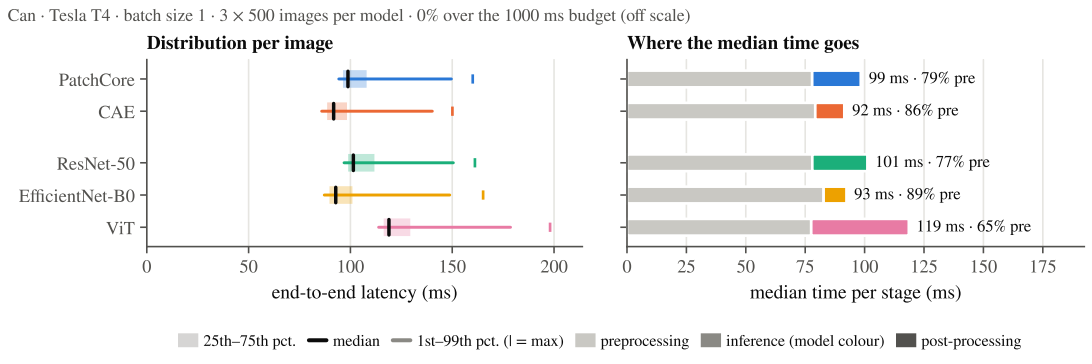
Can · Tesla T4 · batch size 1 · 3 × 500 images per model · 0% over the 1000 ms budget (off scale)
Distribution per image
PatchCore
CAE
ResNet-50
EfficientNet-B0
ViT
0
50
100
150
200
end-to-end latency (ms)
Where the median time goes
99 ms · 79% pre
92 ms · 86% pre
101 ms · 77% pre
93 ms · 89% pre
119 ms · 65% pre
0
25
50
75
100
125
150
175
median time per stage (ms)
25th–75th pct.
median
1st–99th pct. (| = max)
preprocessing
inference (model colour)
post-processing